

DSC540-T302 Data Preparation

Weeks 5 & 6: Data Formats/Data Structures/Data Sources

Saravanan Janarthanan

Assignment

Activity 5.01: Reading Tabular Data from a Web Page and Creating DataFrames

In this activity, you have been given a Wikipedia page where you have the GDP of all countries listed. You have to create three DataFrames from the three sources mentioned on the page ([https://en.wikipedia.org/wiki/List_of_countries_by_GDP_\(nominal\)](https://en.wikipedia.org/wiki/List_of_countries_by_GDP_(nominal))).

You will have to do the following:

Open the page in a separate Chrome/Firefox tab and use something like an Inspect Element tool to view the source HTML and understand its structure.

Read the page using bs4.

Find the table structure you will need to deal with (how many tables are there?).

Find the right table using bs4.

Separate the source names and their corresponding data.

Get the source names from the list of sources you have created.

Separate the header and data from the data that you separated before for the first source only, and then create a DataFrame using that.

Repeat the last task for the other two data sources.

```
In [1]: # Load the modules
from bs4 import BeautifulSoup
import numpy as np
import pandas as pd
```

Read the page using bs4.

Since the page / link has a updated content with only one table, whereas the original html had 3 tables, the backup copy available locally is used

```
In [2]: # Read the page using bs4.

soup = ""
# Open the html file available in local file system and load into a soup object
```

```
with open("datasets/List_of_countries_by_GDP.htm", "r", encoding='utf-8', errors='ignore') as fd:
    soup = BeautifulSoup(fd)
```

```
In [3]: # Find the table structure you will need to deal with (how many tables are there?).
# Find all tables
tables = soup.find_all('table')

print(" Number of tables in the page is : " , len(tables) )
```

Number of tables in the page is : 9

Find the right table using bs4.

Inspecting the table structure the GDP values are listed in 3 tables wrapped by a parent table.

The parent table styled using css class "**wikitable**")}

The 3 tables that hold the GDP values are styles using '**wikitable sortable jquery-tablesorter**' class

```
In [4]: # Retrieve the header by parsing the first 3 rows that has the header for 3 tables using
gdp_tables = soup.find("table", {"class": "wikitable"})
tr_elements = gdp_tables.find_all('tr')
td_elements = tr_elements[0].find_all('td')
tables_hdr = [cell.text.strip() for cell in td_elements]
print(tables_hdr)
```

['Per the International Monetary Fund (2017)[1]', 'Per the World Bank (2017)[20]', 'Per the United Nations (2016)[21][22]']

```
In [5]: # Retrieve all the tags in the table that was styled using class value 'wikitable sortable'
# This returns all the 3 tables that has the gdp values
tables = soup.find_all('table', class_='wikitable sortable jquery-tablesorter')

# Declare a list to hold the 3 tables data in a dataframe for later printing
table_content_df_lst = []

# Extract and print table data
# Iterate through the tables soup object ( 3 tables are available)
for index, table in enumerate(tables, start=1):
    # Retrieve all the TR tags
    rows = table.find_all('tr')

    # Declare a list to hold all the row data of the table
    temp_lst = []

    for row in rows:
        # filter all the TH, TD, THEAD and TBODY tags
        cells = row.find_all(['th', 'td', 'thead', 'tbody'])

        # strip other values and retrieve only the tag text values
        row_data = [cell.text.strip() for cell in cells]

        # based on the web page inspection , two tables has a character and numbers
        # identify the index of the character and assign the values following the character
        spl_car_indx = str(row_data[2]).find("▲")
        if( spl_car_indx != -1):
            row_data[2] = str(row_data[2])[spl_car_indx+1:].strip()
```

```

temp_lst.append(row_data)

# Once all the rows in the tables are process, create a data frame using the list
# the first record is the header and rest are all values
# Add the Dataframe to the list
table_content_df_lst.append(pd.DataFrame(temp_lst[1:], columns=temp_lst[0]))

# Iterate the list and retrieve the dataframe that holds the table data
for idx , df in enumerate(table_content_df_lst):

    # strip and super script reference values in the header
    sub_scr_idx = tables_hdr[idx].find('[')
    print("Table : ", tables_hdr[idx][:sub_scr_idx])

    # print the dataframe top 5 records
    print(df.head())
    print("\n\n")

```

Table : Per the International Monetary Fund (2017)

	Rank	Country	GDP(US\$MM)
0		World[19]	79,865,481
1	1	United States	19,390,600
2	2	China[n 1]	12,014,610
3	3	Japan	4,872,135
4	4	Germany	3,684,816

Table : Per the World Bank (2017)

	Rank	Country	GDP(US\$MM)
0		World	80,683,787
1	1	United States	19,390,604
2		European Union[23]	17,277,698
3	2	China[n 4]	12,237,700
4	3	Japan	4,872,137

Table : Per the United Nations (2016)

	Rank	Country	GDP(US\$MM)
0		World[24]	75,648,448
1	1	United States	18,624,475
2	2	China[n 4]	11,218,281
3	3	Japan	4,936,211
4	4	Germany	3,477,796

Screen shot of the tables in web page

Per the International Monetary Fund (2017) ^[1]			Per the World Bank (2017) ^[2]			Per the United Nations (2016) ^{[2][22]}		
Rank ↕	Country	GDP (US\$MM) ↕	Rank ↕	Country	GDP (US\$MM) ↕	Rank ↕	Country	GDP (US\$MM) ↕
	World^[9]	79,865,481		World	80,683,787		World^[24]	75,648,448
1	United States	19,390,600	1	United States	19,390,604	1	United States	18,624,475
2	China ^[n 1]	12,014,610		European Union ^[23]	17,277,698	2	China ^[n 4]	11,218,281
3	Japan	4,872,135	2	China ^[n 4]	12,237,700	3	Japan	4,936,211
4	Germany	3,684,816	3	Japan	4,872,137	4	Germany	3,477,796
5	United Kingdom	2,624,529	4	Germany	3,677,439	5	United Kingdom	2,647,898
6	India	2,611,012	5	United Kingdom	2,622,434	6	France	2,465,453
7	France	2,583,560	6	India	2,597,491	7	India	2,259,642
8	Brazil	2,054,969	7	France	2,582,501	8	Italy	1,858,913
9	Italy	1,937,894	8	Brazil	2,055,506	9	Brazil	1,795,925
10	Canada	1,652,412	9	Italy	1,934,798	10	Canada	1,529,760
11	South Korea	1,538,030	10	Canada	1,653,043	11	South Korea	1,411,245
12	Russia ^[n 2]	1,527,469	11	Russia ^[n 2]	1,577,524	12	Australia	1,304,463

Activity 6.01: Handling Outliers and Missing Data

In this activity, we will identify and get rid of outliers. Here, we have a CSV file. The goal here is to clean the data by using the knowledge that we have learned about so far and come up with a nicely formatted DataFrame. Identify the type of outliers and their effect on the data and clean the messy data.

The dataset that we have used here can be found in the `visit_data.csv` file. This file contains data generated by a random data generator, and it contains people's names, their gender, `email_id`, `ip_address`, and the number of visits they made to a particular web page.

The steps that will help you solve this activity are as follows:

- Read the `visit_data.csv` file.
- Check for duplicates.
- Check whether any essential column contains NaN.
- Get rid of the outliers.
- Report the size difference.
- Create a box plot to check for outliers.
- Get rid of any outliers

Read the `visit_data.csv` file.

```
In [6]: # Read the visit_data.csv file.
vis_dat_df = pd.read_csv('visit_data.csv')
vis_dat_df.head()
```

```
Out[6]:
```

	id	first_name	last_name	email	gender	ip_address	visit
0	1	Sonny	Dahl	sdahl0@mysql.com	Male	135.36.96.183	1225.0
1	2	NaN	NaN	dhoovart1@hud.gov	NaN	237.165.194.143	919.0
2	3	Gar	Armal	garmal2@technorati.com	NaN	166.43.137.224	271.0
3	4	Chiarra	Nulty	cnulty3@newyorker.com	NaN	139.98.137.108	1002.0
4	5	NaN	NaN	sleaver4@elegantthemes.com	NaN	46.117.117.27	2434.0

```
In [7]: print('The file has ', len(vis_dat_df), ' records')
```

The file has 1000 records

Check for duplicates.

```
In [8]: duplicate_data = vis_dat_df.duplicated()
```

```
In [9]: print('there are ', duplicate_data.sum(), ' duplicate records in the data file ')
```

there are 0 duplicate records in the data file

```
In [10]: # Iterate the columns and print the duplicate values
for col_nm in vis_dat_df.columns:
    print('Total number rows that are duplicated for column ', col_nm, ' is ', vis_c
```

Total number rows that are duplicated for column id is 0
Total number rows that are duplicated for column first_name is 320
Total number rows that are duplicated for column last_name is 299
Total number rows that are duplicated for column email is 0
Total number rows that are duplicated for column gender is 997
Total number rows that are duplicated for column ip_address is 0
Total number rows that are duplicated for column visit is 164

Check whether any essential column contains NaN

```
In [11]: # print the nan counts in each column
vis_dat_df.isna().sum()
```

```
Out[11]: id                0
first_name            296
last_name             296
email                 0
gender                505
ip_address            0
visit                 26
dtype: int64
```

Purpose of the data is to collect the number of visits , if it is null or no data, then it does not serve the purpose, hence removing those outliers

```
In [12]: # drop the records that has NAN values in visits column or feature
vis_data_df_cleaned = vis_dat_df.dropna(subset=['visit'])
```

Report the size difference.

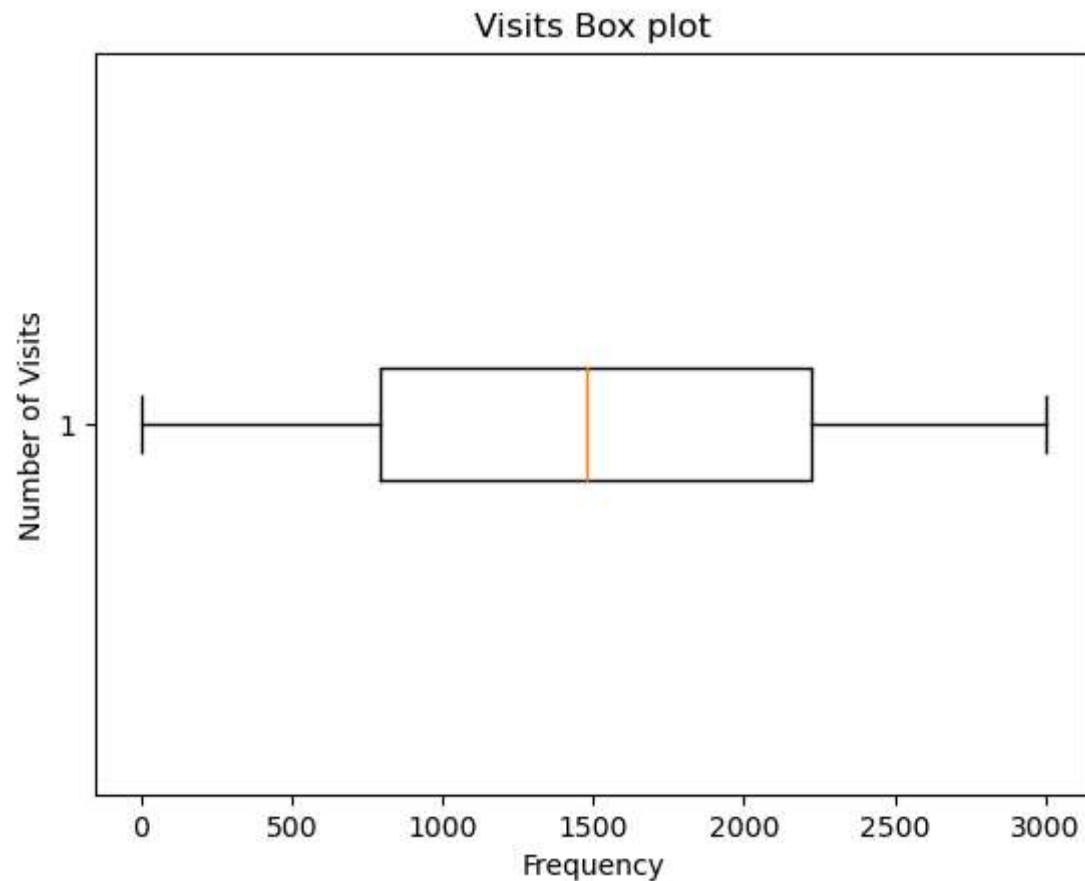
```
In [13]: # Print the net record count after dropping records
print('The key column cleaned up file has ', len(vis_data_df_cleaned), ' records')
print('number of records deleted are ', (1000- len(vis_data_df_cleaned)))
```

The key column cleaned up file has 974 records
number of records deleted are 26

Create a box plot to check for outliers.

```
In [14]: # Create a box plot to visually view the quantiles , Inter Quantile Range, Lower Quan
# and Upper Quantile (right most)
```

```
import matplotlib.pyplot as plt
plt.boxplot(vis_data_df_cleaned['visit'], vert=False, meanline=True)
plt.ylabel('Number of Visits')
plt.xlabel('Frequency')
plt.title('Visits Box plot')
plt.show()
```



```
In [15]: # From the box plot the outliers are on 1st and 4th quantile
# print the quantile values
print('Quantile values')
print(vis_data_df_cleaned['visit'].quantile([ 0,0.25, 0.5, 0.75, 1]))
```

```
Quantile values
0.00    1.00
0.25   794.50
0.50  1477.00
0.75  2224.25
1.00  2998.00
Name: visit, dtype: float64
```

```
In [16]: min(vis_data_df_cleaned['visit'])
```

```
Out[16]: 1.0
```

```
In [17]: max(vis_data_df_cleaned['visit'])
```

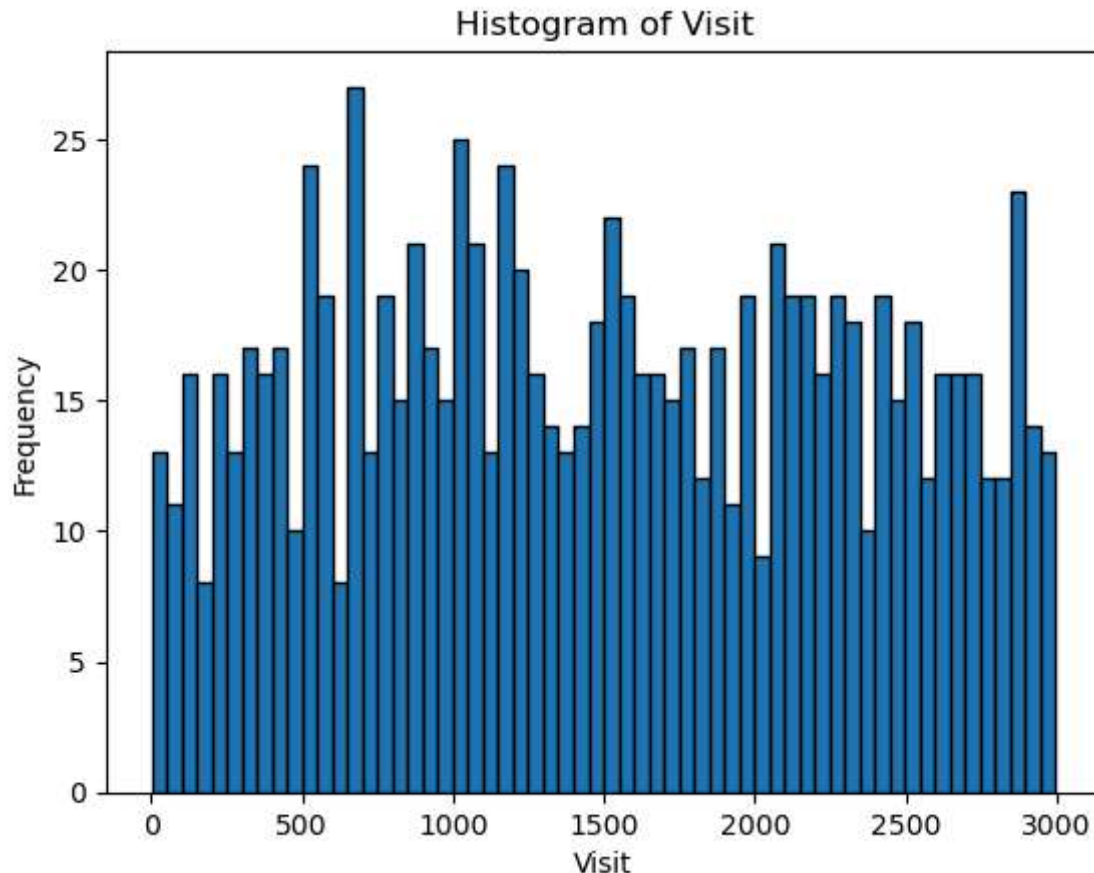
```
Out[17]: 2998.0
```

Get rid of any outliers

```
In [18]: # Create a histogram
plt.hist(vis_data_df_cleaned['visit'], bins=60, edgecolor='black') # Adjust the number of bins

# Add Labels and title
plt.xlabel('Visit')
plt.ylabel('Frequency')
plt.title('Histogram of Visit')

# Show the plot
plt.show()
```



```
In [19]: # From the histogram view
# removing values less than 50 and greater than 2850
vis_data_df_final = vis_data_df_cleaned[(vis_data_df_cleaned['visit'] > 50) & (vis_data_df_cleaned['visit'] < 2850)]
```

```
In [20]: vis_data_df_final.head()
```

```
Out[20]:
```

	id	first_name	last_name	email	gender	ip_address	visit
0	1	Sonny	Dahl	sdahl0@mysql.com	Male	135.36.96.183	1225.0
1	2	NaN	NaN	dhoovart1@hud.gov	NaN	237.165.194.143	919.0
2	3	Gar	Armal	garmal2@technorati.com	NaN	166.43.137.224	271.0
3	4	Chiarra	Nulty	cnulty3@newyorker.com	NaN	139.98.137.108	1002.0
4	5	NaN	NaN	sleaver4@elegantthemes.com	NaN	46.117.117.27	2434.0

```
In [21]: print('Final size of the dataframe after cleaning up the outliers is ', len(vis_data_df_final))
```

Final size of the dataframe after cleaning up the outliers is 913